

Timothy Waldin

San Francisco, CA | timothy@waldin.net | (415) 676-9427
[linkedin.com/in/twaldin](https://www.linkedin.com/in/twaldin) | github.com/twaldin | x.com/twaldin | tim.waldin.net

Software engineer, Reliability @ Lindy.ai — production agent evals and prompt optimization; OSS coding-agent tooling

EXPERIENCE

Software Engineer | Lindy.ai — Reliability Team

May 2026 – Present

- Shipped Online Scoring: evaluation system reading and grading millions of production agent traces per week with LLM judges orchestrated by Temporal workflows
- Implemented GEPA (evolutionary prompt optimization) in production systems handling thousands of user agents, searching for Pareto-optimal prompts over hundreds of evals
- Led Reliability's eval and scoring program: owned the cutover from toy-harness evals to production-faithful replay, plus scoring-honesty fixes after prod incidents so offline numbers match live agent behavior
- Reworked the eval system to wrap production: capture live agent runs as frozen cases and replay them in isolated clones (no live secrets or egress) against current code, with Temporal workers and a cutover so existing suite/score/compare workflows kept working

PROJECTS

hone *TypeScript, Python, GEPA* github.com/twaldin/hone

2025

- GEPA-based prompt optimizer that drives coding-CLI subscriptions through harness as its mutator instead of paid API keys; took Claude Haiku 4.5 from 6/9 to 8/9 solved (65% → 85%) on a held-out GitHub bug set, for ~\$1 in mutator tokens. Being rewritten.

harness *Python, TypeScript* github.com/twaldin/harness

2025

- Extracted the multi-CLI adapter layer into a published Python (`harness-cli`) and TypeScript (`@twaldin/harness-ts`) library covering 13 coding CLIs, with a shared test-fixture suite enforcing identical behavior across both implementations

flt *TypeScript, tmux* github.com/twaldin/flt

2025

- Multi-agent orchestrator — generalized my `tmux-orchestrator` + `claudecord` prototypes to 6 coding CLIs; raw-ANSI TUI with damage-tracked screen buffer, vim keybinds, inter-agent inbox, git-worktree isolation. Published as `@twaldin/flt-cli` on npm

AgentElo *TypeScript, Python* tim.waldin.net/agentelo

2025

- Pairwise Bradley-Terry (Elo-style) leaderboard ranking 148 agents across 6 harnesses on real GitHub PRs — ~1B tokens, \$642 spend, 3.5K verified runs. CLI runs locally to rank your agent against the bundled snapshot

CS2 Trade-Up Bot *TypeScript, React, Express, PostgreSQL, Redis* tradeupbot.app

2025 – Present

- Full-stack market arbitrage platform analyzing profitable CS2 trade-up contracts across 3 marketplaces in real time with 100+ daily active users
- Discovery engine evaluating 100,000+ listing combinations per cycle with Steam OpenID auth, Stripe subscriptions, Cloudflare Turnstile, real-time WebSocket feeds, 17-currency display

term-site *Next.js, Docker, Socket.IO, xterm.js* tim.waldin.net

2025

- Terminal-based portfolio where visitors explore projects inside isolated Docker containers through a browser shell; URL paths deep-link to shell commands (`/blog/<slug>`, `/projects/<name>`)

StudySpot *Next.js, Cloudflare Workers, pgvector, Claude API* studyspot.us

2025

- Co-founded AI study platform using RAG + Claude to answer questions from uploaded course documents or Canvas API import; `pgvector` semantic search, streaming responses, auto-generated flashcards and quizzes

TECHNICAL SKILLS

Languages: TypeScript, JavaScript, Python, C, Shell (Bash, Zsh)

Frameworks: Node.js, React, Next.js, Express, REST APIs, Tailwind, Socket.IO, shadcn/ui

Tools: Git, CI/CD, Docker, Temporal, SQL (PostgreSQL, `pgvector`, SQLite), Redis, nginx, `tmux`, Cloudflare Workers, Bun, uv

AI: agent evals and benchmarking, prompt optimization (GEPA, DSPy), LLM-as-judge pipelines, RAG

EDUCATION

Purdue University | Polytechnic Institute

Aug 2024 – Present

B.S. Web Programming and Design (on leave of absence since May 2026)

Leadership: Founded high school robotics team; competed at FIRST Robotics Competition (FRC) World Championships (2×)